WORKING WITH BIG DATA IN POLITICAL SCIENCE

An introduction to using Google's BigQuery with R

WHAT PROBLEM ARE WE SOLVING?

- Truly big data
  - Need to work outside of your computer's memory
- Need a tool that helps you store, manage, and work with these data
- Don't want to learn a whole new language

BIGQUERY

What is Google's BigQuery?

- Serverless data warehouse
- Built-in query engine

bigrquery

R package that allows you to work with data stored in BigQuery through R.

To install:

```
install.packages(c("bigrquery", "DBI", "dplyr"))
```

To load into your current session:

```
library(bigrquery)
library(DBI)
library(dplyr)
```

HOW TO STORE YOUR DATA

BigQuery is hierarchical:

- **Tables** are stored in:
  - **Datasets**, which are stored in:
    - **Projects**.

HOW TO STORE YOUR DATA

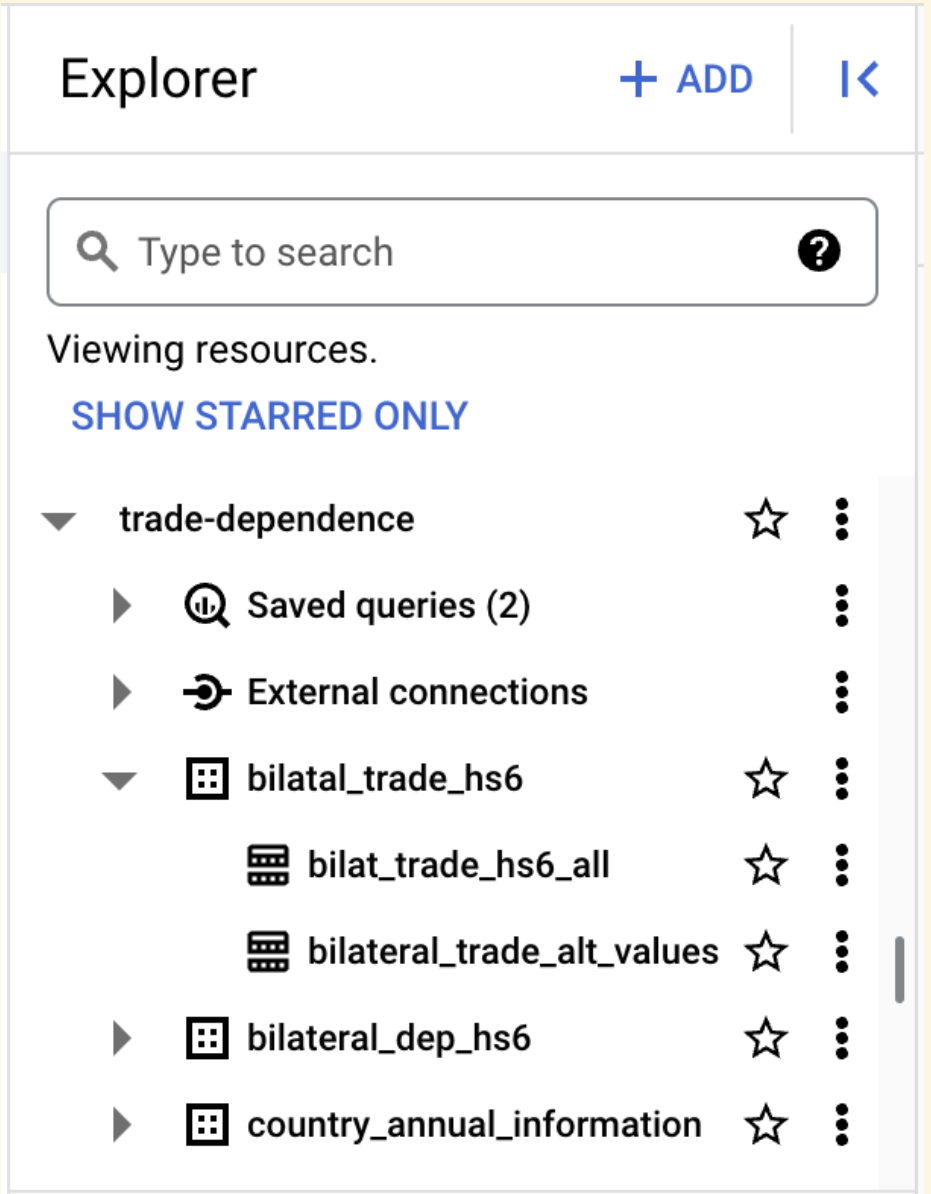

TABLES

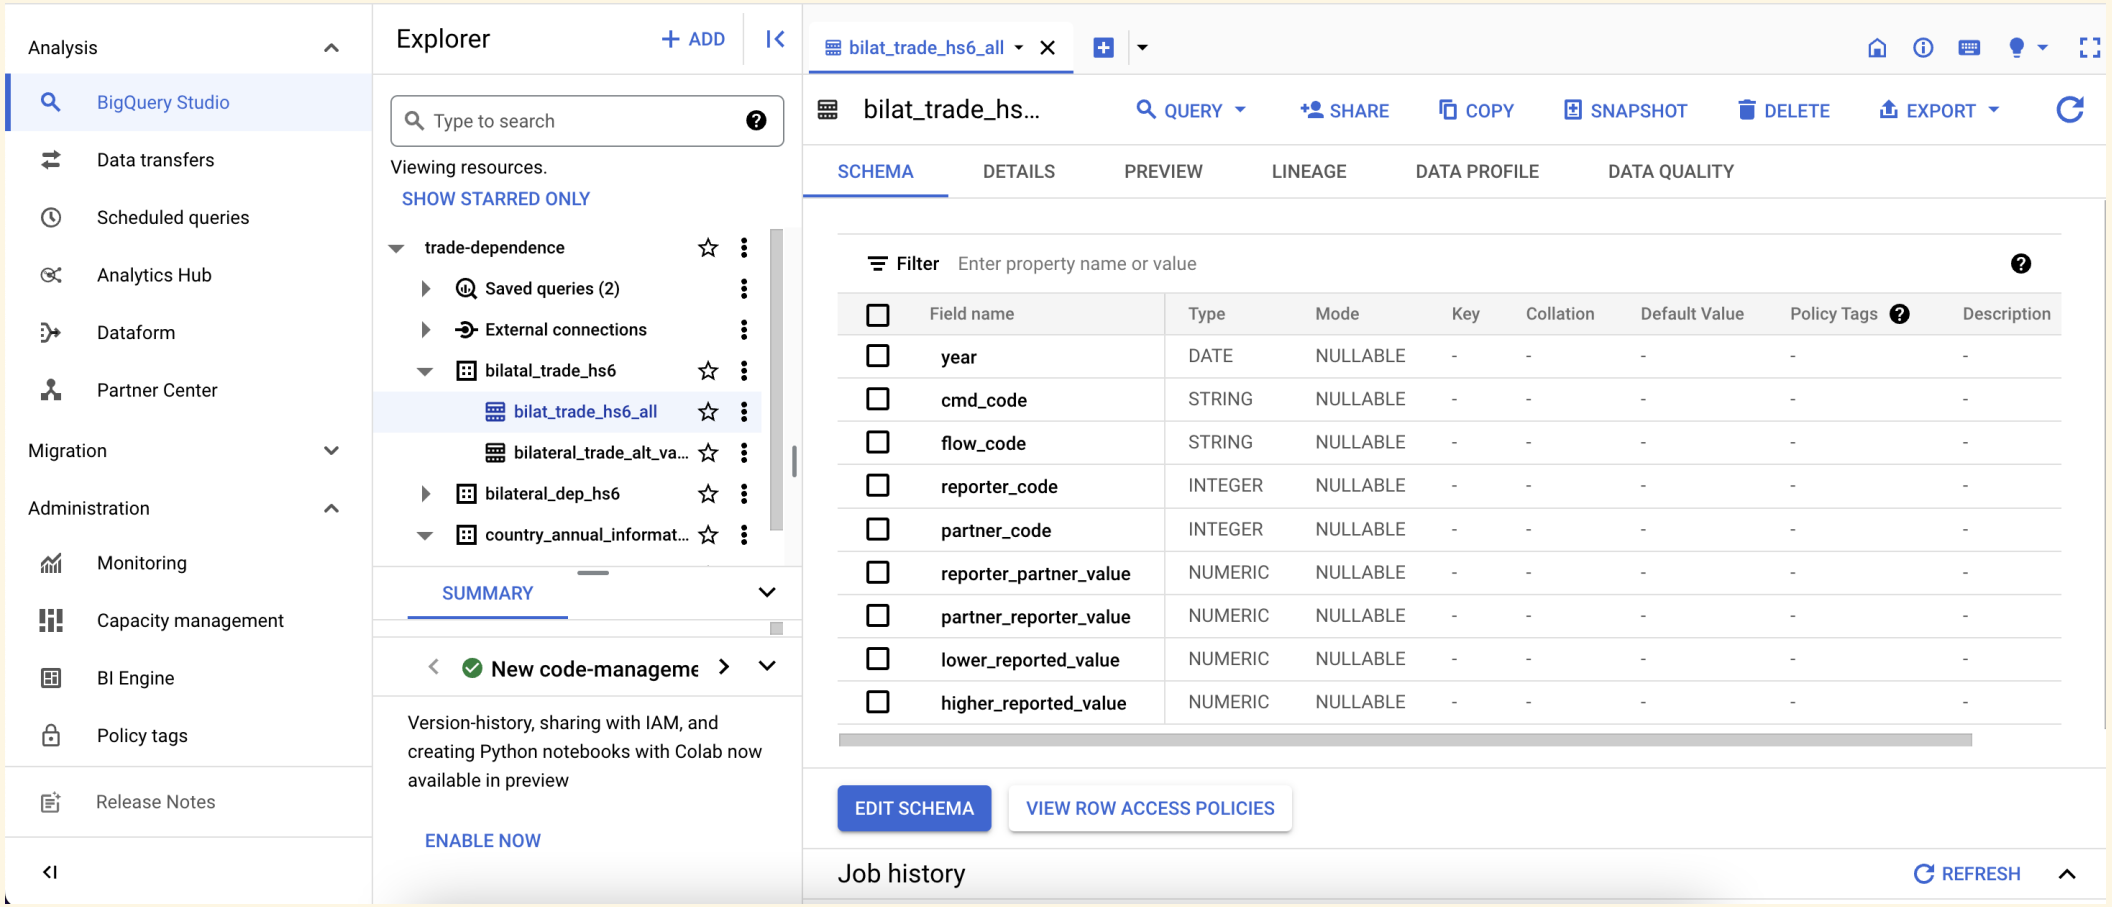
Analysis
BigQuery Studio
Data transfers
Scheduled queries
Analytics Hub
Dataform
Partner Center
Migration
Administration
Monitoring
Capacity management
BI Engine
Policy tags
Release Notes
Explorer
+ ADD
Type to search
Viewing resources.
SHOW STARRED ONLY
trade-dependence
Saved queries (2)
External connections
bilatal_trade_hs6
bilat_trade_hs6_all
bilateral_trade_alt_va...
bilateral_dep_hs6
country_annual_informat...
SUMMARY
New code-manageme
Version-history, sharing with IAM, and creating Python notebooks with Colab now available in preview
ENABLE NOW
bilat_trade_hs6_all
bilat_trade_hs...
QUERY
SHARE
COPY
SNAPSHOT
DELETE
EXPORT
SCHEMA
DETAILS
PREVIEW
LINEAGE
DATA PROFILE
DATA QUALITY
Filter
Enter property name or value
Field name
Type
Mode
Key
Collation
Default Value
Policy Tags
Description
year DATE NULLABLE
cmd_code STRING NULLABLE
flow_code STRING NULLABLE
reporter_code INTEGER NULLABLE
partner_code INTEGER NULLABLE
reporter_partner_value NUMERIC NULLABLE
partner_reporter_value NUMERIC NULLABLE
lower_reported_value NUMERIC NULLABLE
higher_reported_value NUMERIC NULLABLE
EDIT SCHEMA
VIEW ROW ACCESS POLICIES
Job history
REFRESH

TABLES

bilat_trade_hs... QUERY SHARE COPY SNAPSHOT DELETE EXPORT

SCHEMA DETAILS PREVIEW LINEAGE DATA PROFILE DATA QUALITY

Partitioned on field	year
Partition expiration	Partitions do not expire
Partition filter	Not required

Storage info

Number of rows	283,167,024
Number of partitions	16
Total logical bytes	23.06 GB
Active logical bytes	0 B
Long term logical bytes	23.06 GB
Total physical bytes	5.31 GB
Active physical bytes	0 B
Long term physical bytes	5.31 GB
Time travel physical bytes	0 B

TABLES

bilat_trade_hs... QUERY SHARE COPY SNAPSHOT DELETE EXPORT

SCHEMA DETAILS PREVIEW LINEAGE DATA PROFILE DATA QUALITY

Row	year	cmd_code	flow_code	reporter_code	partner_code	reporter_partner	partn
1	2020-01-01	999999	M	4	32	null	98
2	2020-01-01	240220	M	4	51	null	33
3	2020-01-01	240399	M	4	51	null	
4	2020-01-01	854449	M	4	36	null	9
5	2020-01-01	180632	M	4	36	null	
6	2020-01-01	300510	M	4	36	null	
7	2020-01-01	300650	M	4	36	null	
8	2020-01-01	300230	M	4	36	null	
9	2020-01-01	870899	M	4	36	null	
10	2020-01-01	950699	M	4	36	null	
11	2020-01-01	630790	M	4	36	null	
12	2020-01-01	630590	M	4	36	null	
13	2020-01-01	902519	M	4	36	null	
14	2020-01-01	902140	M	4	36	null	
15	2020-01-01	392330	M	4	36	null	

Results per page: 50 1 – 50 of 283167024

UPLOADING YOUR DATA

We will step through uploading your data to BigQuery

UPLOADING DATA MANUALLY

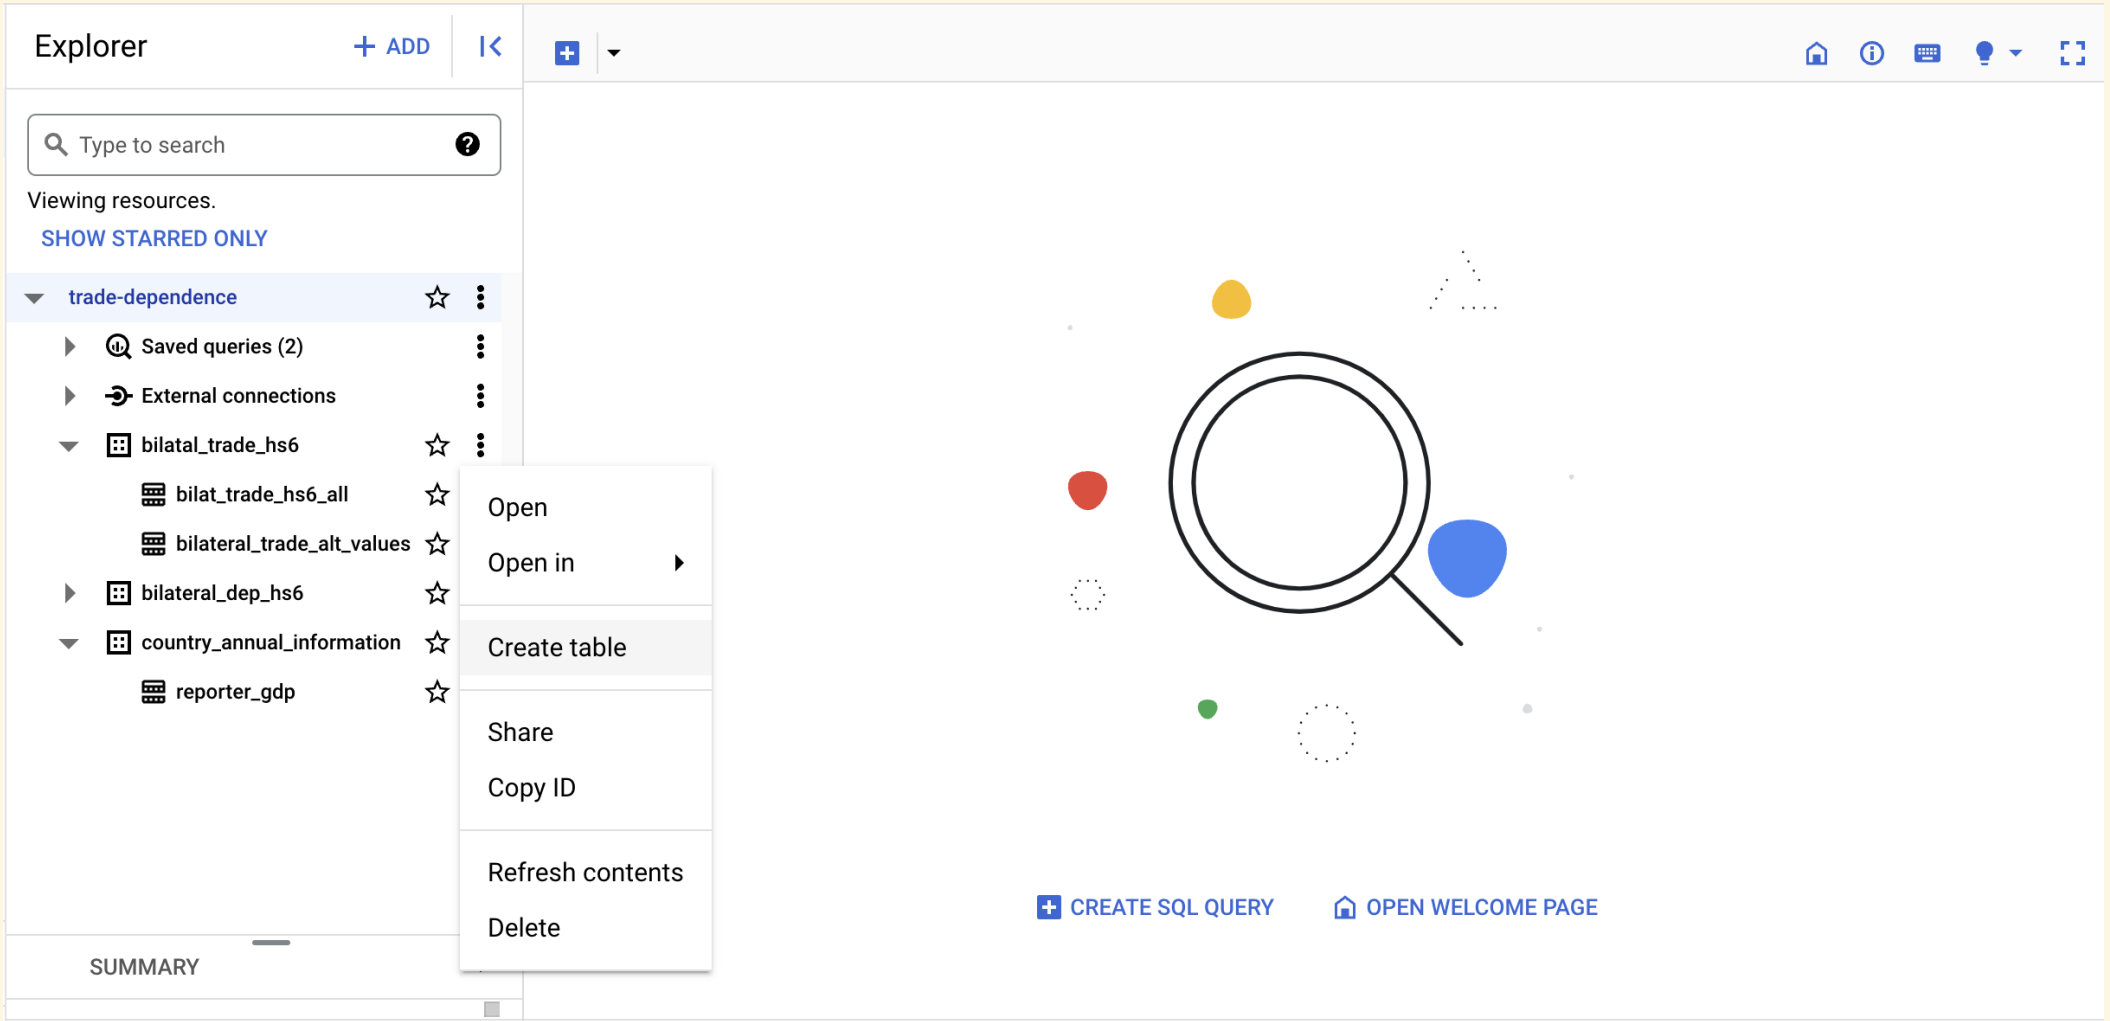

UPLOADING DATA MANUALLY

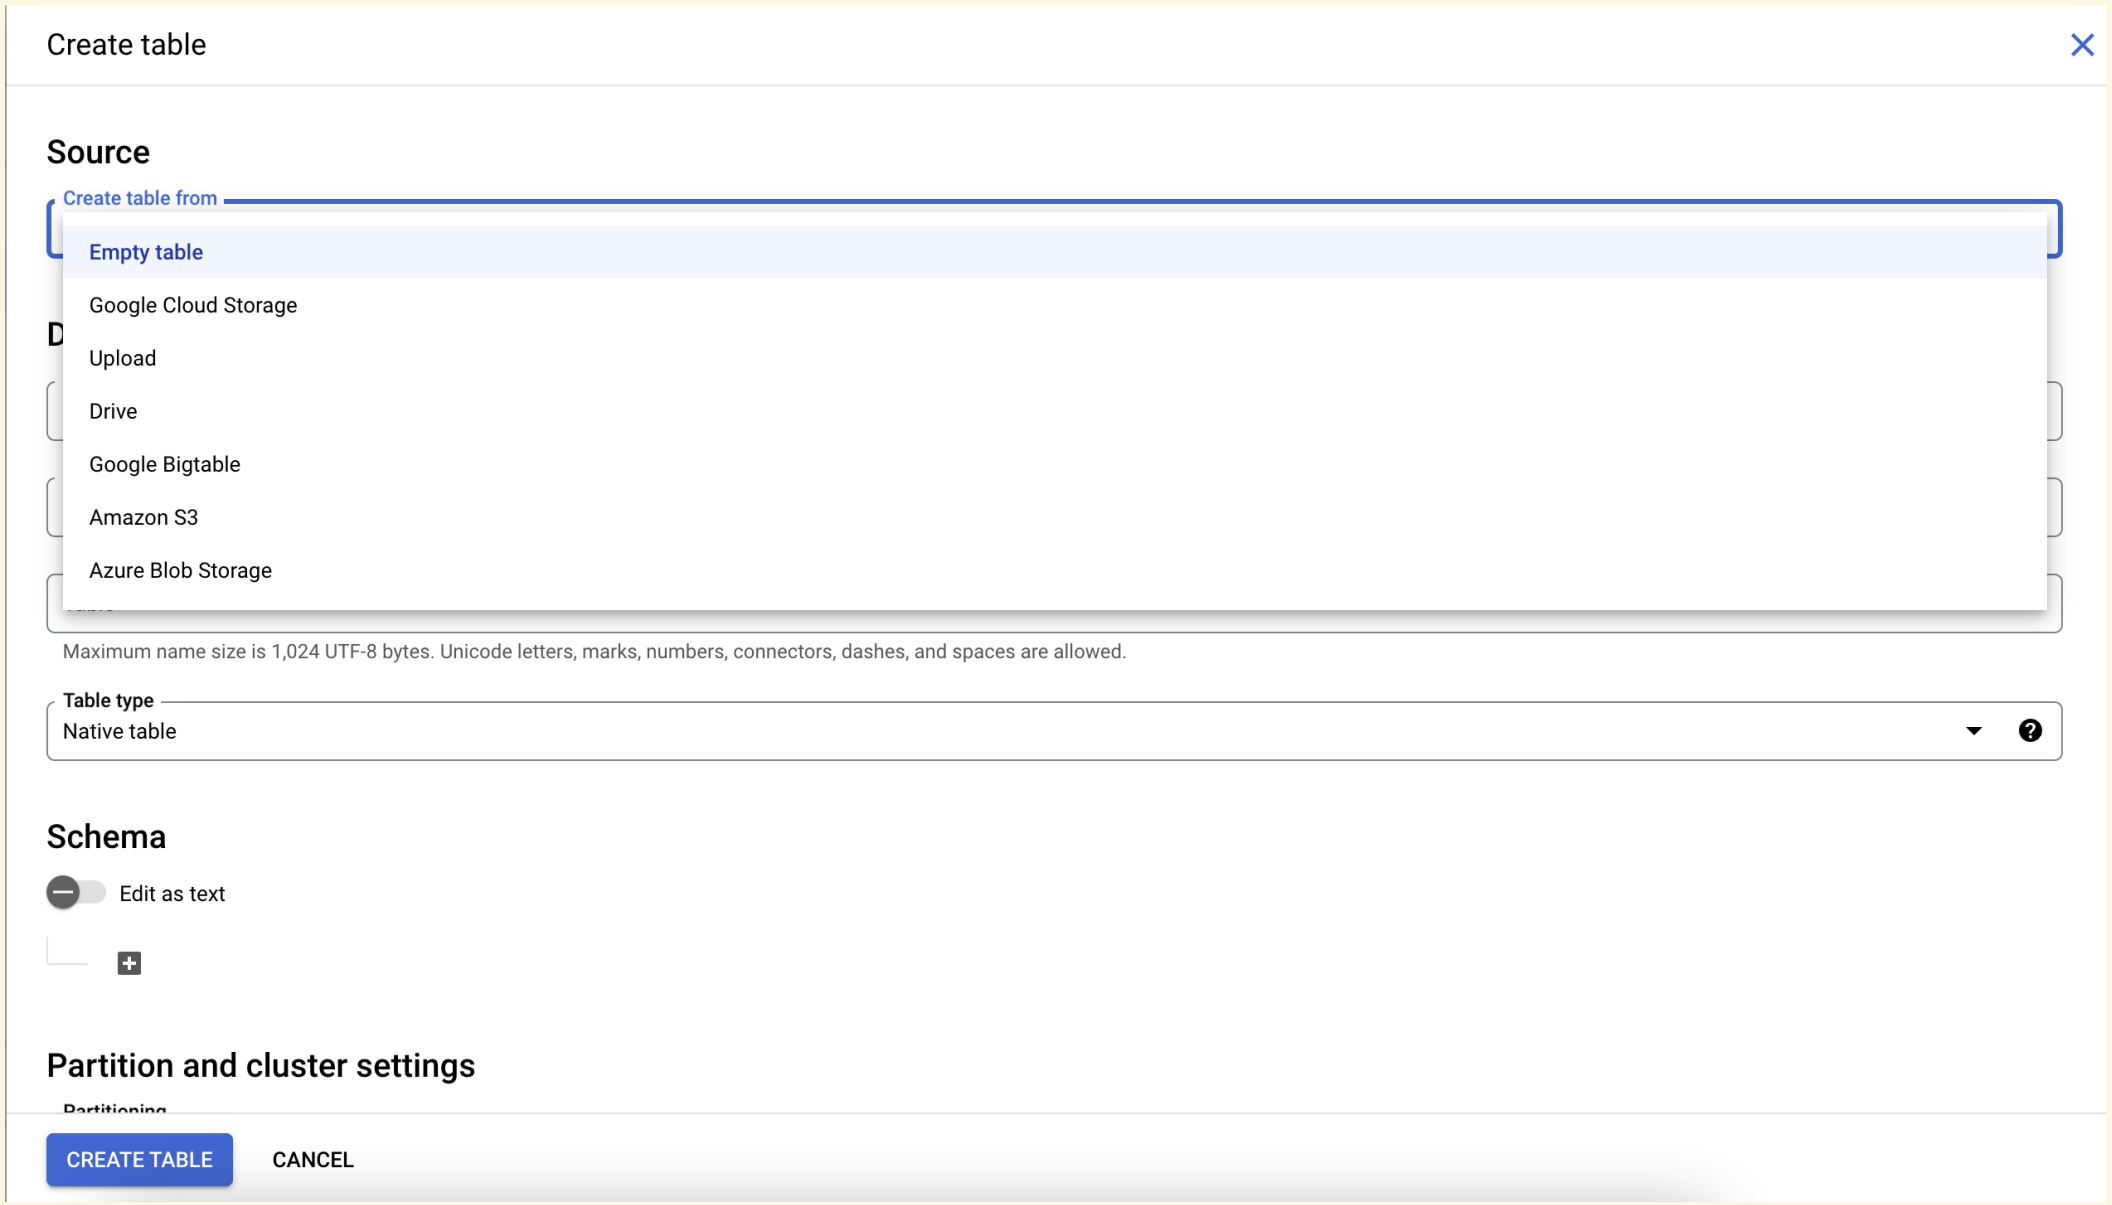

UPLOADING DATA MANUALLY

Create table

Source

Create table from: Upload

Select file *: API_NY.GDP.MKTP.CD_DS2_en_csv_v2_125.csv BROWSE

File format: CSV

Destination

Project *: trade-dependence BROWSE

Dataset *: bilatal_trade_hs6

Table *: GDP_current

Maximum name size is 1,024 UTF-8 bytes. Unicode letters, marks, numbers, connectors, dashes, and spaces are allowed.

Table type: Native table

Schema

Auto detect

CREATE TABLE CANCEL

UPLOADING DATA MANUALLY

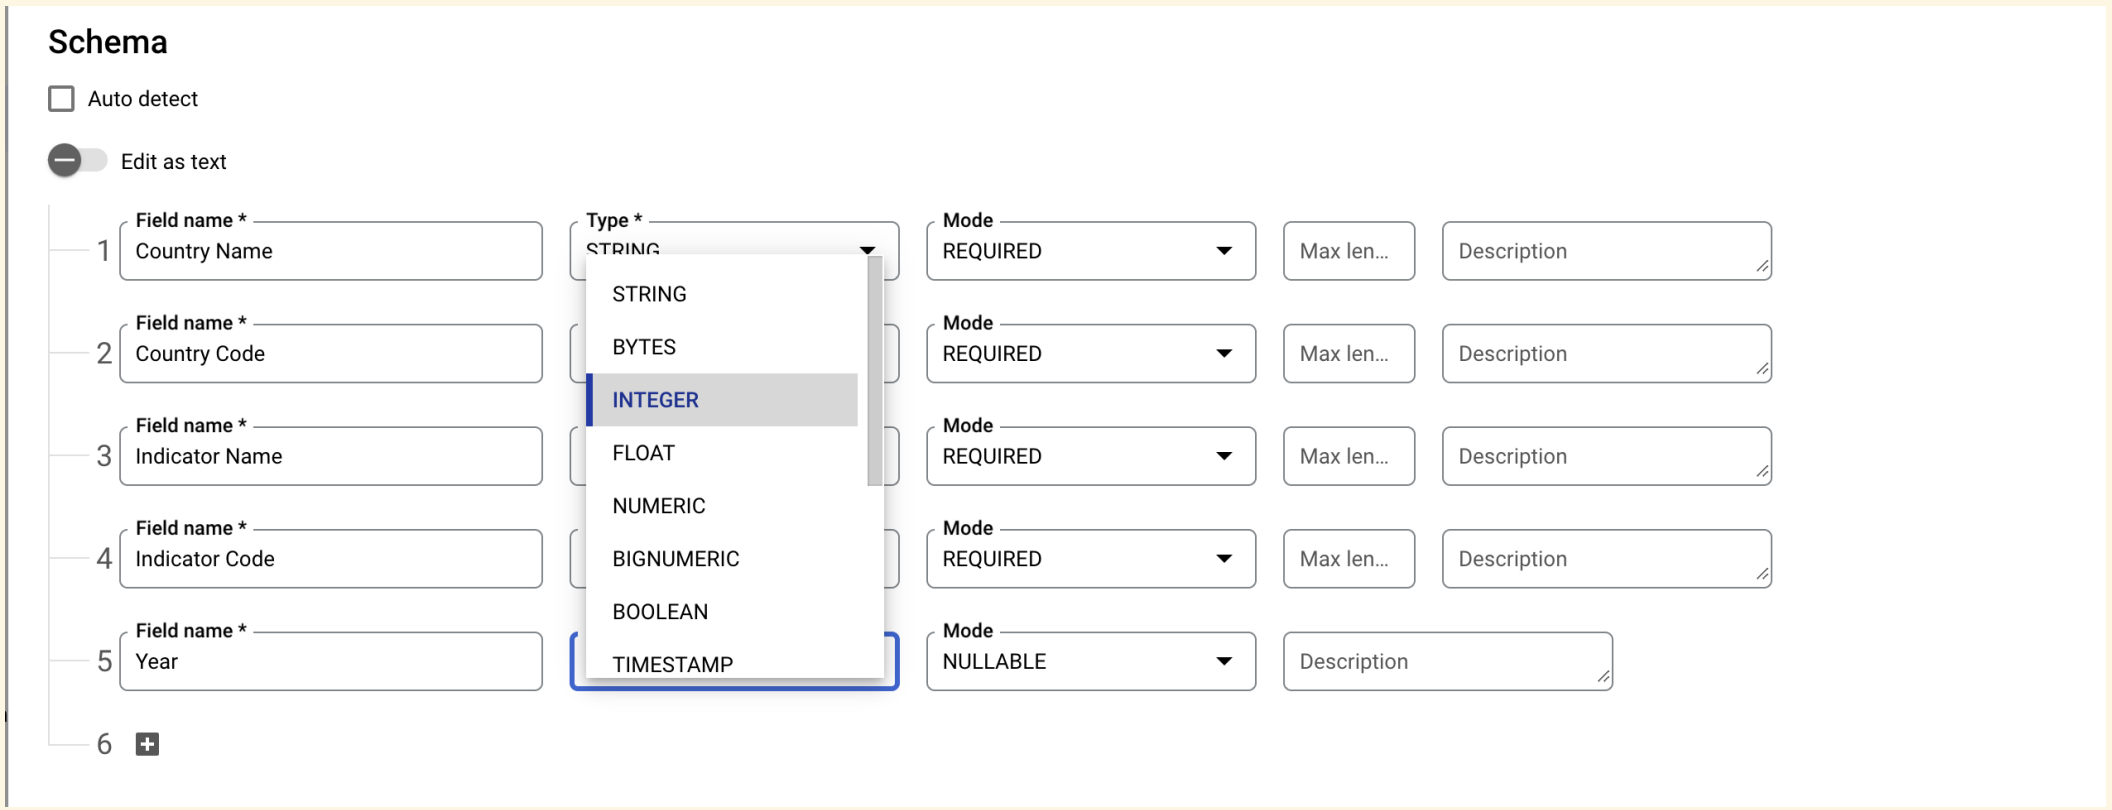

UPLOADING DATA MANUALLY

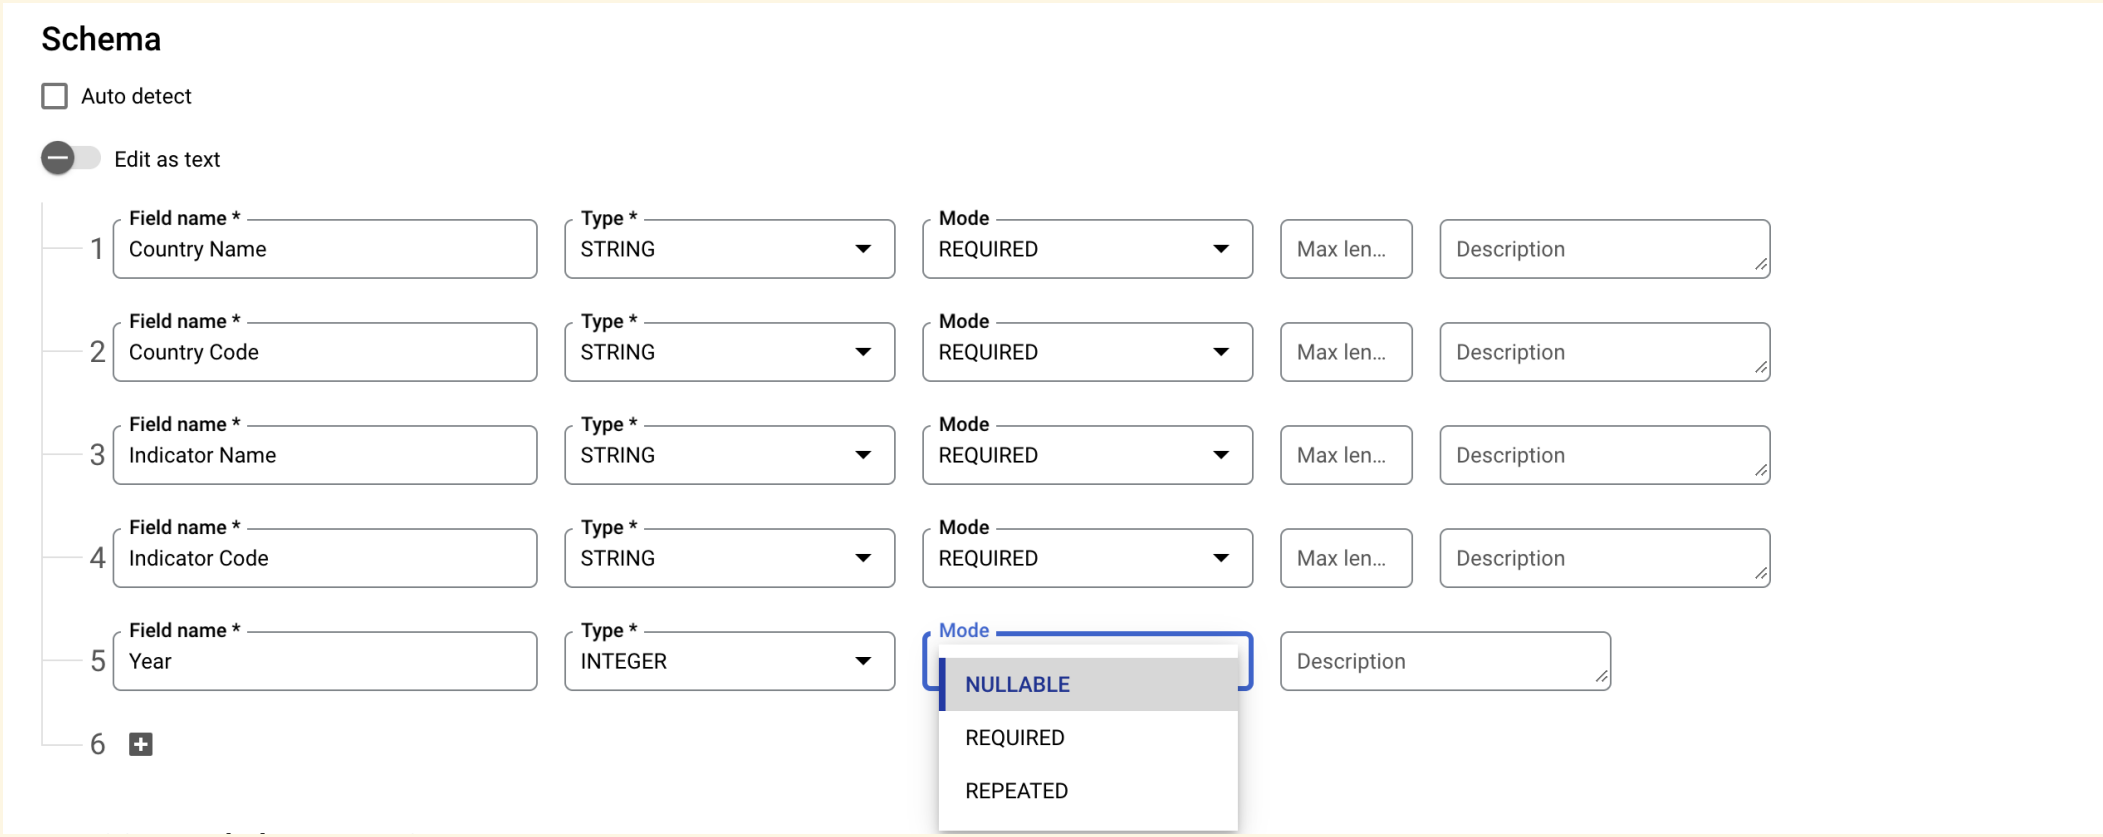

UPLOADING DATA USING `bigrquery`

```
library(bigrquery)
library(DBI)
library(dplyr)
```

List useful information:

```
selected_project <- "trade-dependence"
selected_dataset <- "bilatal_trade_hs6"
```

UPLOADING DATA USING `bigrquery`

First, check that the table does not already exist:

```
bq_gdp_current <- bq_table(project = selected_project,
                           dataset = selected_dataset,
                           table = "gdp_current")

bq_table_exists(bq_gdp_current)
```

```
[1] FALSE
```

The first time you do this, you will need to authorize `bigrquery`'s access to your

UPLOADING DATA USING `bigrquery`

Next, create the (empty) table:

```
bq_table_create(
  bq_gdp_current,
  fields = gdp_current,
  friendly_name = "GDP (current USD)",
  description = "The data was extracted from the World Bank."
)

bq_table_exists(bq_gdp_current)
```

```
[1] TRUE
```

UPLOADING DATA USING `bigrquery`

Next, upload your data to that empty table:

```
1 bq_table_upload(bq_gdp_current, gdp_df)
```

WORKING WITH YOUR DATA

We will now step through how to work with big data out of your computer's memory

STARTING IN R WITH `dplyr`

Let's collect data on Australia's GDP.

STARTING IN R WITH `dplyr`

These data are stored in the `trade-dependence` project and the `country_annual_information` dataset.

```
selected_project <- "trade-dependence"
selected_dataset <- "country_annual_information"

con <- dbConnect(
  bigrquery::bigquery(),
  project = selected_project,
  dataset = selected_dataset,
  billing = selected_project
)

con
```

```
<BigQueryConnection>
  Dataset: trade-dependence.country_annual_information
  Billing: trade-dependence
```

STARTING IN R WITH dplyr

Create the connection to the reporter_gdp table:

```
1 gdp_df <- tbl(con, "reporter_gdp")
2 gdp_df
```

```
# Source:   table<reporter_gdp> [?? x 3]
# Database: BigQueryConnection
   year       reporter_code reporter_gdp_current
   <date>             <int>                <dbl>
 1 2003-01-01            92                   NA
 2 2003-01-01           136                   NA
 3 2003-01-01           531                   NA
 4 2003-01-01           292                   NA
 5 2003-01-01           408                   NA
 6 2003-01-01            NA                   NA
 7 2003-01-01           520                   NA
 8 2003-01-01           534                   NA
 9 2003-01-01           706                   NA
10 2003-01-01           728                   NA
# ℹ more rows
```

STARTING IN R WITH dplyr

Query that table:

```r
aus_gdp <- gdp_df |>
  filter(reporter_code == 36) |>
  collect()

aus_gdp
```

```
# A tibble: 23 × 3
   year       reporter_code reporter_gdp_current
   <date>             <dbl>                <dbl>
 1 2002-01-01            36             3.96e11
 2 2018-01-01            36             1.43e12
 3 2009-01-01            36             9.29e11
 4 2011-01-01            36             1.40e12
 5 2004-01-01            36             6.14e11
 6 2021-01-01            36             1.55e12
 7 2017-01-01            36             1.33e12
 8 2008-01-01            36             1.06e12
 9 2001-01-01            36             3.79e11
10 2012-01-01            36             1.55e12
# ℹ 13 more rows
```

MOVING BETWEEN R AND BIGQUERY

R will write your SQL queries for you:

```r
gdp_df |>
  filter(reporter_code == 36) |>
  show_query()
```

```
<SQL>
SELECT `reporter_gdp`.*
FROM `reporter_gdp`
WHERE (`reporter_code` = 36.0)
```

MOVING BETWEEN R AND BIGQUERY

You can perform that query in BigQuery's in-built query engine:

MOVING BETWEEN R AND BIGQUERY

You can perform that query in BigQuery's in-built query engine:

PERFORMING COMPLEX QUERIES

munge_alt_trade_value RUN SAVE QUERY SHARE SCHEDULE MORE This query will process 905.8 MB when run.

```
INSERT INTO trade-dependence.bilatal_trade_hs6.bilateral_trade_alt_values
SELECT
  `year`,
  `cmd_code`,
  `flow_code`,
  `reporter_code`,
  `partner_code`,
  SUM(`alt_reporter_partner_value`) AS `alt_reporter_partner_value`,
  SUM(`alt_partner_reporter_value`) AS `alt_partner_reporter_value`,
  SUM(`alt_lower_reported_value`) AS `alt_lower_reported_value`,
  SUM(`alt_higher_reported_value`) AS `alt_higher_reported_value`
FROM (
  SELECT
    `LHS`.`year` AS `year`,
    `LHS`.`cmd_code` AS `cmd_code`,
    `LHS`.`flow_code` AS `flow_code`,
    `reporter_code`,
    `partner_code`,
    `alt_partner_code`,
    `alt_partner_partner_code`,
    `alt_reporter_partner_value`,
    `alt_partner_reporter_value`,
    `alt_lower_reported_value`,
    `alt_higher_reported_value`
  FROM (
    SELECT *
    FROM (
      SELECT `year`, `cmd_code`, `flow_code`, `reporter_code`, `partner_code`
      FROM `trade-dependence.bilatal_trade_hs6.bilat_trade_hs6_all`
    ) `q01`
    WHERE (`reporter_code` != `partner_code`) AND (`year` = '2005-01-01')
  ) `LHS`
  LEFT JOIN (
    SELECT
      `year`,
      `cmd_code`,
      CASE WHEN (`flow_code` = 'M') THEN 'X' WHEN NOT (`flow_code` = 'M') THEN 'M' END AS `flow_code`,
```

NEXT STEPS

- Partitioning and clustering your datasets
  - Great for yearly, country-level data
- Integrated ML model-building